

K Means

April 24, 2019

This project illustrates the use of KMeans Clustering to cluster Universities data into to two groups, Private and Public.

I have used a data frame with 777 observations on the following 18 variables.

Private A factor with levels No and Yes indicating private or public university

Apps Number of applications received

Accept Number of applications accepted

Enroll Number of new students enrolled

Top10perc Pct. new students from top 10% of H.S. class

Top25perc Pct. new students from top 25% of H.S. class

F.Undergrad Number of fulltime undergraduates

P.Undergrad Number of parttime undergraduates

Outstate Out-of-state tuition

Room.Board Room and board costs

Books Estimated book costs

Personal Estimated personal spending

PhD Pct. of faculty with Ph.D.'s

Terminal Pct. of faculty with terminal degree

S.F.Ratio Student/faculty ratio

perc.alumni Pct. alumni who donate

Expend Instructional expenditure per student

Grad.Rate Graduation rate

```
In [ ]: import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import numpy as np
```

```
In [38]: df = pd.read_csv('College_Data',)
```

```
In [39]: df.head()
```

```
Out[39]:
```

	Private	Apps	Accept	Enroll	Top10perc	\
Abilene Christian University	Yes	1660	1232	721	23	
Adelphi University	Yes	2186	1924	512	16	
Adrian College	Yes	1428	1097	336	22	
Agnes Scott College	Yes	417	349	137	60	
Alaska Pacific University	Yes	193	146	55	16	

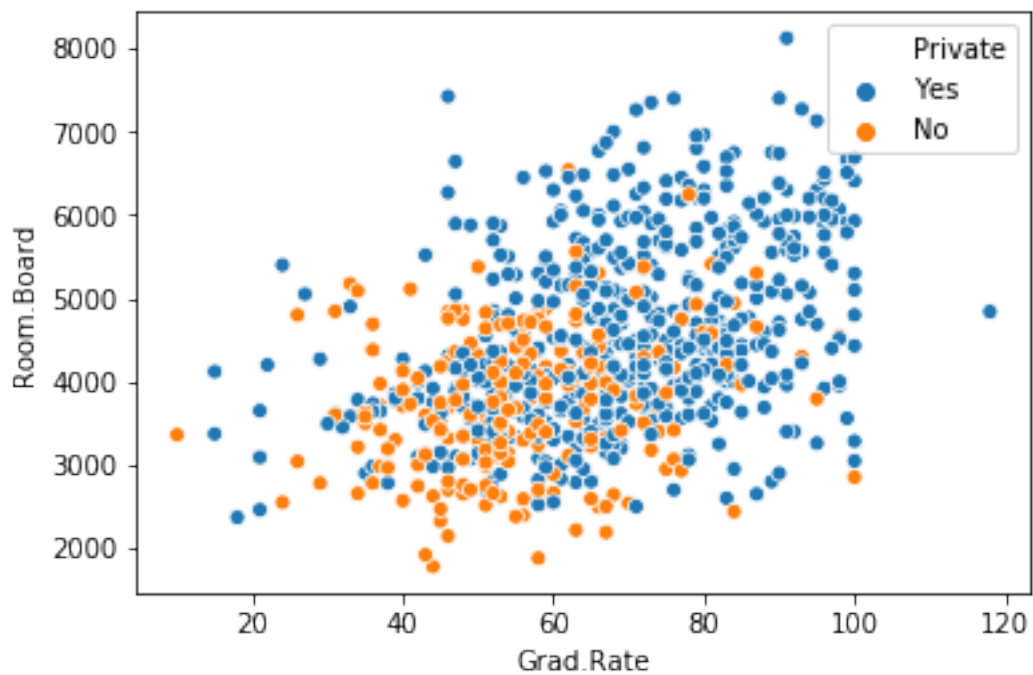
	Top25perc	F.Undergrad	P.Undergrad	Outstate	\
Abilene Christian University	52	2885	537	7440	
Adelphi University	29	2683	1227	12280	
Adrian College	50	1036	99	11250	
Agnes Scott College	89	510	63	12960	
Alaska Pacific University	44	249	869	7560	

	Room.Board	Books	Personal	PhD	Terminal	\
Abilene Christian University	3300	450	2200	70	78	
Adelphi University	6450	750	1500	29	30	
Adrian College	3750	400	1165	53	66	
Agnes Scott College	5450	450	875	92	97	
Alaska Pacific University	4120	800	1500	76	72	

	S.F.Ratio	perc.alumni	Expend	Grad.Rate
Abilene Christian University	18.1	12	7041	60
Adelphi University	12.2	16	10527	56
Adrian College	12.9	30	8735	54
Agnes Scott College	7.7	37	19016	59
Alaska Pacific University	11.9	2	10922	15

```
In [40]: sns.scatterplot(df['Grad.Rate'],df['Room.Board'],hue=df['Private'])
```


```
Out[40]: <matplotlib.axes._subplots.AxesSubplot at 0x15e13ad8748>
```



```
In [41]: sns.distplot(df['Grad.Rate'],bins = 6)
```

```
C:\ProgramData\Anaconda3\lib\site-packages\scipy\stats\stats.py:1713: FutureWarning: Using a non-tuple sequence for multidimensional indexing is deprecated; using `tuple` instead. Errors may arise from you using this construct with arrays that are writable in the future.
return np.add.reduce(sorted[indexer] * weights, axis=axis) / sumval
```

```
Out[41]: <matplotlib.axes._subplots.AxesSubplot at 0x15e13ad7b38>
```



From the above plots, we can see that there is record with grad rate > 100, which is not possible in general, so we will be setting the value to 100 using pandas functions.

```
In [42]: df.describe()
```

```
Out[42]:
```

	Apps	Accept	Enroll	Top10perc	Top25perc	\
count	777.000000	777.000000	777.000000	777.000000	777.000000	
mean	3001.638353	2018.804376	779.972973	27.558559	55.796654	
std	3870.201484	2451.113971	929.176190	17.640364	19.804778	
min	81.000000	72.000000	35.000000	1.000000	9.000000	
25%	776.000000	604.000000	242.000000	15.000000	41.000000	
50%	1558.000000	1110.000000	434.000000	23.000000	54.000000	
75%	3624.000000	2424.000000	902.000000	35.000000	69.000000	
max	48094.000000	26330.000000	6392.000000	96.000000	100.000000	

	F.Undergrad	P.Undergrad	Outstate	Room.Board	Books	\
count	777.000000	777.000000	777.000000	777.000000	777.000000	
mean	3699.907336	855.298584	10440.669241	4357.526384	549.380952	

std	4850.420531	1522.431887	4023.016484	1096.696416	165.105360
min	139.000000	1.000000	2340.000000	1780.000000	96.000000
25%	992.000000	95.000000	7320.000000	3597.000000	470.000000
50%	1707.000000	353.000000	9990.000000	4200.000000	500.000000
75%	4005.000000	967.000000	12925.000000	5050.000000	600.000000
max	31643.000000	21836.000000	21700.000000	8124.000000	2340.000000

	Personal	PhD	Terminal	S.F.Ratio	perc.alumni	\
count	777.000000	777.000000	777.000000	777.000000	777.000000	
mean	1340.642214	72.660232	79.702703	14.089704	22.743887	
std	677.071454	16.328155	14.722359	3.958349	12.391801	
min	250.000000	8.000000	24.000000	2.500000	0.000000	
25%	850.000000	62.000000	71.000000	11.500000	13.000000	
50%	1200.000000	75.000000	82.000000	13.600000	21.000000	
75%	1700.000000	85.000000	92.000000	16.500000	31.000000	
max	6800.000000	103.000000	100.000000	39.800000	64.000000	

	Expend	Grad.Rate
count	777.000000	777.000000
mean	9660.171171	65.46332
std	5221.768440	17.17771
min	3186.000000	10.00000
25%	6751.000000	53.00000
50%	8377.000000	65.00000
75%	10830.000000	78.00000
max	56233.000000	118.00000

```
In [43]: df[df['Grad.Rate']>100]['Grad.Rate']
```

```
Out[43]: Cazenovia College      118
          Name: Grad.Rate, dtype: int64
```

```
In [44]: df['Grad.Rate'].replace(118,100,inplace=True)
```

Importing K means from scikit learn and fitting the dataframe to divide the dataset into two clusters using kmeans.

```
In [45]: from sklearn.cluster import KMeans
```

```
In [46]: kmeans = KMeans(n_clusters = 2)
```

```
In [48]: kmeans.fit(df.drop('Private',axis = 1))
```

```
Out[48]: KMeans(algorithm='auto', copy_x=True, init='k-means++', max_iter=300,
                n_clusters=2, n_init=10, n_jobs=None, precompute_distances='auto',
                random_state=None, tol=0.0001, verbose=0)
```

```
In [49]: kmeans.cluster_centers_
```

```
Out [49]: array([[1.81323468e+03, 1.28716592e+03, 4.91044843e+02, 2.53094170e+01,
                5.34708520e+01, 2.18854858e+03, 5.95458894e+02, 1.03957085e+04,
                4.31136472e+03, 5.41982063e+02, 1.28033632e+03, 7.04424514e+01,
                7.78251121e+01, 1.40997010e+01, 2.31748879e+01, 8.93204634e+03,
                6.50926756e+01],
                [1.03631389e+04, 6.55089815e+03, 2.56972222e+03, 4.14907407e+01,
                7.02037037e+01, 1.30619352e+04, 2.46486111e+03, 1.07191759e+04,
                4.64347222e+03, 5.95212963e+02, 1.71420370e+03, 8.63981481e+01,
                9.13333333e+01, 1.40277778e+01, 2.00740741e+01, 1.41705000e+04,
                6.75925926e+01]])
```

Have further evaluated the model to check for accuracy as we have private variable already present in the dataset. This is not usually done in Kmeans, but doing it just because we have data available.

```
In [50]: def converter(cluster):
        if cluster=='Yes':
            return 1
        else:
            return 0
```

```
In [51]: df['Cluster'] = df['Private'].apply(converter)
```

```
In [52]: df.head()
```

```
Out [52]:
```

	Private	Apps	Accept	Enroll	Top10perc	\
Abilene Christian University	Yes	1660	1232	721	23	
Adelphi University	Yes	2186	1924	512	16	
Adrian College	Yes	1428	1097	336	22	
Agnes Scott College	Yes	417	349	137	60	
Alaska Pacific University	Yes	193	146	55	16	

	Top25perc	F.Undergrad	P.Undergrad	Outstate	\
Abilene Christian University	52	2885	537	7440	
Adelphi University	29	2683	1227	12280	
Adrian College	50	1036	99	11250	
Agnes Scott College	89	510	63	12960	
Alaska Pacific University	44	249	869	7560	

	Room.Board	Books	Personal	PhD	Terminal	\
Abilene Christian University	3300	450	2200	70	78	
Adelphi University	6450	750	1500	29	30	
Adrian College	3750	400	1165	53	66	
Agnes Scott College	5450	450	875	92	97	
Alaska Pacific University	4120	800	1500	76	72	

	S.F.Ratio	perc.alumni	Expend	Grad.Rate	\
Abilene Christian University	18.1	12	7041	60	
Adelphi University	12.2	16	10527	56	

Adrian College	12.9	30	8735	54
Agnes Scott College	7.7	37	19016	59
Alaska Pacific University	11.9	2	10922	15

	Cluster
Abilene Christian University	1
Adelphi University	1
Adrian College	1
Agnes Scott College	1
Alaska Pacific University	1

```
In [53]: from sklearn.metrics import confusion_matrix,classification_report
print(confusion_matrix(df['Cluster'],kmeans.labels_))
print(classification_report(df['Cluster'],kmeans.labels_))
```

```
[[138  74]
 [531  34]]
```

	precision	recall	f1-score	support
0	0.21	0.65	0.31	212
1	0.31	0.06	0.10	565
micro avg	0.22	0.22	0.22	777
macro avg	0.26	0.36	0.21	777
weighted avg	0.29	0.22	0.16	777

```
In [ ]:
```