

PCA Project

April 24, 2019

1 Principal Component Analysis

In this project, I have implemented PCA to reduce a 30 dimensions dataset into 2 Principal components and then checked how logistic regression would work with 2 variables and compared it to a 30 features model.

Importing the required libraries

```
In [1]: import matplotlib.pyplot as plt
import pandas as pd
import numpy as np
import seaborn as sns
%matplotlib inline
```

I have used Cancer data set that is available online for doing this.

```
In [17]: df=pd.read_csv('data.csv')
```

```
In [18]: df.head()
```

```
Out[18]:
```

	id	diagnosis	radius_mean	texture_mean	perimeter_mean	area_mean	\
0	842302	M	17.99	10.38	122.80	1001.0	
1	842517	M	20.57	17.77	132.90	1326.0	
2	84300903	M	19.69	21.25	130.00	1203.0	
3	84348301	M	11.42	20.38	77.58	386.1	
4	84358402	M	20.29	14.34	135.10	1297.0	

	smoothness_mean	compactness_mean	concavity_mean	concave points_mean	\
0	0.11840	0.27760	0.3001	0.14710	
1	0.08474	0.07864	0.0869	0.07017	
2	0.10960	0.15990	0.1974	0.12790	
3	0.14250	0.28390	0.2414	0.10520	
4	0.10030	0.13280	0.1980	0.10430	

	...	texture_worst	perimeter_worst	area_worst	smoothness_worst	\
0	...	17.33	184.60	2019.0	0.1622	
1	...	23.41	158.80	1956.0	0.1238	
2	...	25.53	152.50	1709.0	0.1444	

3	...	26.50	98.87	567.7	0.2098
4	...	16.67	152.20	1575.0	0.1374

	compactness_worst	concavity_worst	concave points_worst	symmetry_worst	\
0	0.6656	0.7119	0.2654	0.4601	
1	0.1866	0.2416	0.1860	0.2750	
2	0.4245	0.4504	0.2430	0.3613	
3	0.8663	0.6869	0.2575	0.6638	
4	0.2050	0.4000	0.1625	0.2364	

	fractal_dimension_worst	Unnamed: 32
0	0.11890	NaN
1	0.08902	NaN
2	0.08758	NaN
3	0.17300	NaN
4	0.07678	NaN

[5 rows x 33 columns]

```
In [19]: df.pop('id')
df.head()
```

```
Out[19]:
```

	diagnosis	radius_mean	texture_mean	perimeter_mean	area_mean	\
0	M	17.99	10.38	122.80	1001.0	
1	M	20.57	17.77	132.90	1326.0	
2	M	19.69	21.25	130.00	1203.0	
3	M	11.42	20.38	77.58	386.1	
4	M	20.29	14.34	135.10	1297.0	

	smoothness_mean	compactness_mean	concavity_mean	concave points_mean	\
0	0.11840	0.27760	0.3001	0.14710	
1	0.08474	0.07864	0.0869	0.07017	
2	0.10960	0.15990	0.1974	0.12790	
3	0.14250	0.28390	0.2414	0.10520	
4	0.10030	0.13280	0.1980	0.10430	

	symmetry_mean	...	texture_worst	perimeter_worst	area_worst	\
0	0.2419	...	17.33	184.60	2019.0	
1	0.1812	...	23.41	158.80	1956.0	
2	0.2069	...	25.53	152.50	1709.0	
3	0.2597	...	26.50	98.87	567.7	
4	0.1809	...	16.67	152.20	1575.0	

	smoothness_worst	compactness_worst	concavity_worst	concave points_worst	\
0	0.1622	0.6656	0.7119	0.2654	
1	0.1238	0.1866	0.2416	0.1860	
2	0.1444	0.4245	0.4504	0.2430	
3	0.2098	0.8663	0.6869	0.2575	

4	0.1374	0.2050	0.4000	0.1625
---	--------	--------	--------	--------

	symmetry_worst	fractal_dimension_worst	Unnamed: 32
0	0.4601	0.11890	NaN
1	0.2750	0.08902	NaN
2	0.3613	0.08758	NaN
3	0.6638	0.17300	NaN
4	0.2364	0.07678	NaN

[5 rows x 32 columns]

```
In [21]: df['diagnosis'].replace('M',1,inplace=True)
df['diagnosis'].replace('B',0,inplace=True)
```

```
In [22]: df.head()
```

```
Out[22]:
```

	diagnosis	radius_mean	texture_mean	perimeter_mean	area_mean	\
0	1	17.99	10.38	122.80	1001.0	
1	1	20.57	17.77	132.90	1326.0	
2	1	19.69	21.25	130.00	1203.0	
3	1	11.42	20.38	77.58	386.1	
4	1	20.29	14.34	135.10	1297.0	

	smoothness_mean	compactness_mean	concavity_mean	concave points_mean	\
0	0.11840	0.27760	0.3001	0.14710	
1	0.08474	0.07864	0.0869	0.07017	
2	0.10960	0.15990	0.1974	0.12790	
3	0.14250	0.28390	0.2414	0.10520	
4	0.10030	0.13280	0.1980	0.10430	

	symmetry_mean	...	texture_worst	perimeter_worst	area_worst	\
0	0.2419	...	17.33	184.60	2019.0	
1	0.1812	...	23.41	158.80	1956.0	
2	0.2069	...	25.53	152.50	1709.0	
3	0.2597	...	26.50	98.87	567.7	
4	0.1809	...	16.67	152.20	1575.0	

	smoothness_worst	compactness_worst	concavity_worst	concave points_worst	\
0	0.1622	0.6656	0.7119	0.2654	
1	0.1238	0.1866	0.2416	0.1860	
2	0.1444	0.4245	0.4504	0.2430	
3	0.2098	0.8663	0.6869	0.2575	
4	0.1374	0.2050	0.4000	0.1625	

	symmetry_worst	fractal_dimension_worst	Unnamed: 32
0	0.4601	0.11890	NaN
1	0.2750	0.08902	NaN
2	0.3613	0.08758	NaN

3	0.6638	0.17300	NaN
4	0.2364	0.07678	NaN

[5 rows x 32 columns]

In [23]: df_x = df.drop('diagnosis',axis=1)

In [33]: df_x.head()

```
Out [33]:
```

	radius_mean	texture_mean	perimeter_mean	area_mean	smoothness_mean	\
0	17.99	10.38	122.80	1001.0	0.11840	
1	20.57	17.77	132.90	1326.0	0.08474	
2	19.69	21.25	130.00	1203.0	0.10960	
3	11.42	20.38	77.58	386.1	0.14250	
4	20.29	14.34	135.10	1297.0	0.10030	

	compactness_mean	concavity_mean	concave points_mean	symmetry_mean	\
0	0.27760	0.3001	0.14710	0.2419	
1	0.07864	0.0869	0.07017	0.1812	
2	0.15990	0.1974	0.12790	0.2069	
3	0.28390	0.2414	0.10520	0.2597	
4	0.13280	0.1980	0.10430	0.1809	

	fractal_dimension_mean	...	texture_worst	perimeter_worst	\
0	0.07871	...	17.33	184.60	
1	0.05667	...	23.41	158.80	
2	0.05999	...	25.53	152.50	
3	0.09744	...	26.50	98.87	
4	0.05883	...	16.67	152.20	

	area_worst	smoothness_worst	compactness_worst	concavity_worst	\
0	2019.0	0.1622	0.6656	0.7119	
1	1956.0	0.1238	0.1866	0.2416	
2	1709.0	0.1444	0.4245	0.4504	
3	567.7	0.2098	0.8663	0.6869	
4	1575.0	0.1374	0.2050	0.4000	

	concave points_worst	symmetry_worst	fractal_dimension_worst	Unnamed: 32
0	0.2654	0.4601	0.11890	NaN
1	0.1860	0.2750	0.08902	NaN
2	0.2430	0.3613	0.08758	NaN
3	0.2575	0.6638	0.17300	NaN
4	0.1625	0.2364	0.07678	NaN

[5 rows x 31 columns]

In [25]: df_y = df['diagnosis']

Below steps are performed to standardize the independent variables before using PCA.

```
In [26]: from sklearn.preprocessing import StandardScaler
```

```
In [27]: scaler = StandardScaler()  
        scaler.fit(df_x)
```

```
C:\ProgramData\Anaconda3\lib\site-packages\sklearn\utils\extmath.py:776: RuntimeWarning: invalid  
  updated_mean = (last_sum + new_sum) / updated_sample_count  
C:\ProgramData\Anaconda3\lib\site-packages\sklearn\utils\extmath.py:781: RuntimeWarning: Degree  
  new_unnormalized_variance = np.nanvar(X, axis=0) * new_sample_count
```

```
Out[27]: StandardScaler(copy=True, with_mean=True, with_std=True)
```

```
In [28]: scaled_data = scaler.transform(df_x)
```

```
In [31]: df_xnew = pd.DataFrame(data = scaled_data[0:,0:],columns=df_x.columns)
```

```
In [34]: df_xnew.pop('Unnamed: 32')  
        df_xnew.head()
```

```
Out[34]:
```

	radius_mean	texture_mean	perimeter_mean	area_mean	smoothness_mean	\
0	1.097064	-2.073335	1.269934	0.984375	1.568466	
1	1.829821	-0.353632	1.685955	1.908708	-0.826962	
2	1.579888	0.456187	1.566503	1.558884	0.942210	
3	-0.768909	0.253732	-0.592687	-0.764464	3.283553	
4	1.750297	-1.151816	1.776573	1.826229	0.280372	

	compactness_mean	concavity_mean	concave	points_mean	symmetry_mean	\
0	3.283515	2.652874		2.532475	2.217515	
1	-0.487072	-0.023846		0.548144	0.001392	
2	1.052926	1.363478		2.037231	0.939685	
3	3.402909	1.915897		1.451707	2.867383	
4	0.539340	1.371011		1.428493	-0.009560	

	fractal_dimension_mean	...	radius_worst	\
0	2.255747	...	1.886690	
1	-0.868652	...	1.805927	
2	-0.398008	...	1.511870	
3	4.910919	...	-0.281464	
4	-0.562450	...	1.298575	

	texture_worst	perimeter_worst	area_worst	smoothness_worst	\
0	-1.359293	2.303601	2.001237	1.307686	
1	-0.369203	1.535126	1.890489	-0.375612	
2	-0.023974	1.347475	1.456285	0.527407	
3	0.133984	-0.249939	-0.550021	3.394275	
4	-1.466770	1.338539	1.220724	0.220556	

	compactness_worst	concavity_worst	concave	points_worst	symmetry_worst	\
--	-------------------	-----------------	---------	--------------	----------------	---

0	2.616665	2.109526	2.296076	2.750622
1	-0.430444	-0.146749	1.087084	-0.243890
2	1.082932	0.854974	1.955000	1.152255
3	3.893397	1.989588	2.175786	6.046041
4	-0.313395	0.613179	0.729259	-0.868353

	fractal_dimension_worst
0	1.937015
1	0.281190
2	0.201391
3	4.935010
4	-0.397100

[5 rows x 30 columns]

I have used PCA from scikit learn and specified no of components to be 2 and created 2 Principal components.

```
In [35]: from sklearn.decomposition import PCA
```

```
In [36]: pca = PCA(n_components=2)
```

```
In [37]: pca.fit(df_xnew)
```

```
Out[37]: PCA(copy=True, iterated_power='auto', n_components=2, random_state=None,
          svd_solver='auto', tol=0.0, whiten=False)
```

Now we can transform this data to its first 2 principal components.

```
In [38]: x_pca = pca.transform(df_xnew)
```

```
In [39]: df_xnew.shape
```

```
Out[39]: (569, 30)
```


```
In [40]: x_pca.shape
```

```
Out[40]: (569, 2)
```

Plotting the 2 Principal components

```
In [42]: plt.figure(figsize=(8,6))
          plt.scatter(x_pca[:,0],x_pca[:,1],c=df_y,cmap='plasma')
          plt.xlabel('First principal component')
          plt.ylabel('Second Principal Component')
```

```
Out[42]: Text(0, 0.5, 'Second Principal Component')
```



Clearly by using these two components we can easily separate these two classes.

The components correspond to combinations of the original features, the components themselves are stored as an attribute of the fitted PCA object:

In [43]: `pca.components_`


```
Out[43]: array([[ 0.21890244,  0.10372458,  0.22753729,  0.22099499,  0.14258969,
                  0.23928535,  0.25840048,  0.26085376,  0.13816696,  0.06436335,
                  0.20597878,  0.01742803,  0.21132592,  0.20286964,  0.01453145,
                  0.17039345,  0.15358979,  0.1834174 ,  0.04249842,  0.10256832,
                  0.22799663,  0.10446933,  0.23663968,  0.22487053,  0.12795256,
                  0.21009588,  0.22876753,  0.25088597,  0.12290456,  0.13178394],
                 [-0.23385713, -0.05970609, -0.21518136, -0.23107671,  0.18611302,
                  0.15189161,  0.06016536, -0.0347675 ,  0.19034877,  0.36657547,
                  -0.10555215,  0.08997968, -0.08945723, -0.15229263,  0.20443045,
                  0.2327159 ,  0.19720728,  0.13032156,  0.183848 ,  0.28009203,
                  -0.21986638, -0.0454673 , -0.19987843, -0.21935186,  0.17230435,
                  0.14359317,  0.09796411, -0.00825724,  0.14188335,  0.27533947]])
```

In this numpy matrix array, each row represents a principal component, and each column relates back to the original features. we can visualize this relationship with a heatmap:

In [44]: `df_comp = pd.DataFrame(pca.components_, columns=df_xnew.columns)`

```
In [45]: plt.figure(figsize=(12,6))
         sns.heatmap(df_comp,cmap='plasma',)
```

```
Out[45]: <matplotlib.axes._subplots.AxesSubplot at 0x2f01c2e3630>
```



This heatmap shows correlation between each component and actual variables

```
In [48]: df_pca = pd.DataFrame(data=x_pca[0:,0:],columns=['Component_1', 'Component_2'])
```

```
In [49]: df_pca.head()
```

```
Out[49]:
```

	Component_1	Component_2
0	9.192837	1.948583
1	2.387802	-3.768172
2	5.733896	-1.075174
3	7.122953	10.275589
4	3.935302	-1.948072

Doing Logistic Regression using Actual Variables

```
In [50]: from sklearn.model_selection import train_test_split
```

```
In [51]: x1_train,x1_test,y1_train,y1_test = train_test_split(df_xnew,df_y,test_size=0.3,random
```



```
In [52]: from sklearn.linear_model import LogisticRegression
```

```
In [53]: logmodel=LogisticRegression()
```

```
In [54]: logmodel.fit(x1_train,y1_train)
```

```
C:\ProgramData\Anaconda3\lib\site-packages\sklearn\linear_model\logistic.py:433: FutureWarning  
FutureWarning)
```

```
Out[54]: LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=True,  
    intercept_scaling=1, max_iter=100, multi_class='warn',  
    n_jobs=None, penalty='l2', random_state=None, solver='warn',  
    tol=0.0001, verbose=0, warm_start=False)
```

```
In [55]: pred_y1 = logmodel.predict(x1_test)
```

```
In [56]: from sklearn.metrics import classification_report,confusion_matrix
```

```
In [58]: print('Before PCA:')  
    print('\n')  
    print(confusion_matrix(y1_test,pred_y1))  
    print('\n')  
    print(classification_report(y1_test,pred_y1))
```

Before PCA:

```
[[104  0]  
 [  3 64]]
```

	precision	recall	f1-score	support
0	0.97	1.00	0.99	104
1	1.00	0.96	0.98	67
micro avg	0.98	0.98	0.98	171
macro avg	0.99	0.98	0.98	171
weighted avg	0.98	0.98	0.98	171

```
In [60]: x1_train.shape
```

```
Out[60]: (398, 30)
```

Doing Logistic Regression Using PCA components

```
In [61]: x2_train,x2_test,y2_train,y2_test = train_test_split(df_pca,df_y,test_size=0.3,random
```

```
In [62]: from sklearn.linear_model import LogisticRegression
```

```
In [63]: logmodel2=LogisticRegression()
```

```
In [64]: logmodel2.fit(x2_train,y2_train)
```

```
C:\ProgramData\Anaconda3\lib\site-packages\sklearn\linear_model\logistic.py:433: FutureWarning
FutureWarning)
```

```
Out[64]: LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=True,
intercept_scaling=1, max_iter=100, multi_class='warn',
n_jobs=None, penalty='l2', random_state=None, solver='warn',
tol=0.0001, verbose=0, warm_start=False)
```

```
In [65]: pred_y2 = logmodel2.predict(x2_test)
```

```
In [66]: from sklearn.metrics import classification_report,confusion_matrix
```

```
In [69]: print('After PCA:')
print('\n')
print(confusion_matrix(y2_test,pred_y2))
print('\n')
print(classification_report(y2_test,pred_y2))
```

After PCA:

```
[[103  1]
 [ 9 58]]
```

	precision	recall	f1-score	support
0	0.92	0.99	0.95	104
1	0.98	0.87	0.92	67
micro avg	0.94	0.94	0.94	171
macro avg	0.95	0.93	0.94	171
weighted avg	0.94	0.94	0.94	171

From the results of logistic regression. we can see that we had an accuracy of 0.98 with 30 dimensions, and when reduced to 2 dimensions using PCA, we still had 0.94 dimensions. This shows the power of PCA in rapidly reducing the dimensions.

```
In [ ]:
```